

2-3: Ownership Rules and Borrowing (Theory)

Artem Pavlov, TII, Abu Dhabi, 15.05.2025

Stack, heap, and static memory

- The stack stores values in the order it gets them and removes the values in the opposite order (i.e. LIFO)
 - Very efficient
 - Has limited size (usually 2-8 MiB)
 - Does not need any runtime support
 - Size of objects should be known at compile time
- The heap stores objects of arbitrary size and its memory is managed by an “allocator”
 - Less efficient
 - Requires existence of an “allocator”
 - Can store dynamically sized objects
- Static memory is allocated at compile time and often is part of generated binary

Ownership

- Ownership is a set of rules that govern how a Rust program manages memory.
- Ownership rules:
 - Each value in Rust has an owner.
 - There can only be one owner at a time.
 - When the owner goes out of scope, the value will be dropped.

Scopes

```
{           // s is not valid here, it's not yet declared
  let s = "hello"; // s is valid from this point forward

  // do stuff with s
}           // this scope is now over, and s is no longer valid

{
  let mut s = String::from("hello");
  s.push_str(", world!"); // appends a literal to a String
  println!("{}", s);
  // s gets dropped here
}
```

Move and Copy

```
#[derive(Debug)]
struct Foo(u32);
let x = Foo(42);
let y = x;


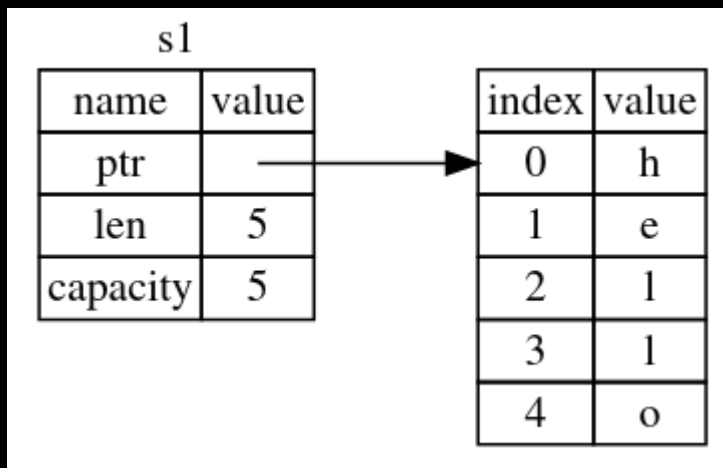
println!("{y:?}");
// x can not be used
// println!("{x:?}");

let x = 42u32;
let y = x;
// both x and y can be used
println!("{x:?} {y:?}");
```

- Moving value performs bitwise copy
- Move usually invalidates old location
- If type implements the **Copy** trait, then old location stays valid

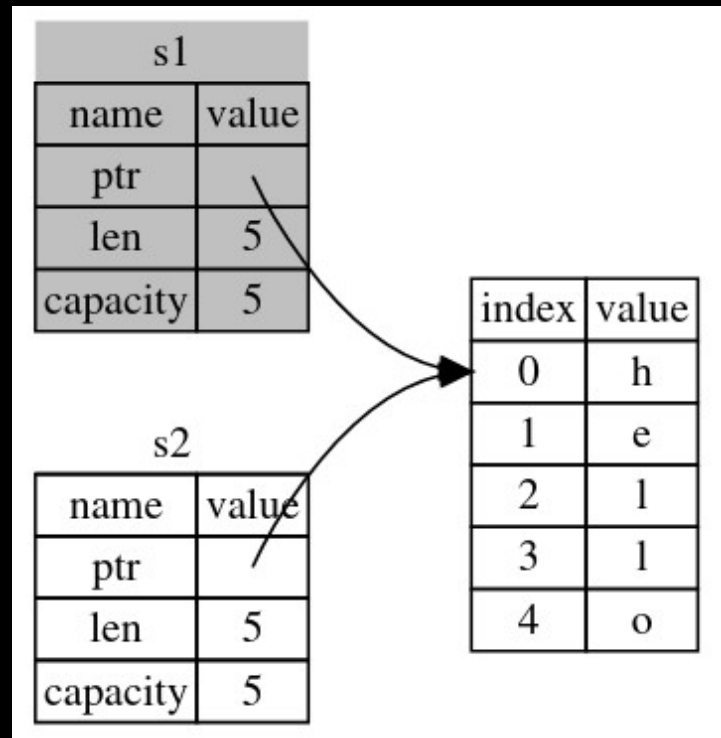
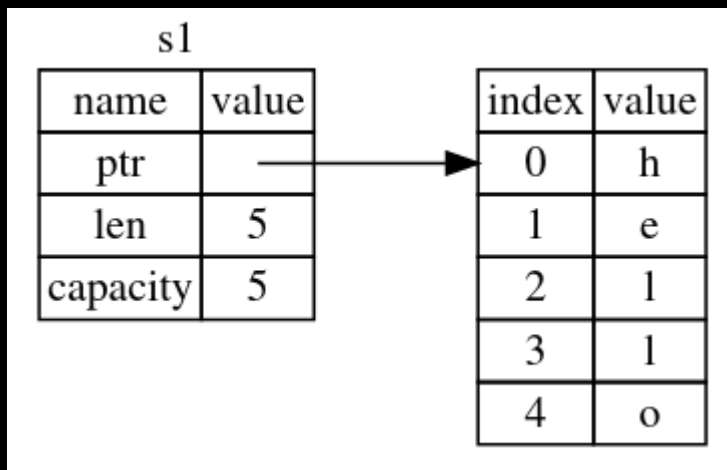
Move and heap-allocated data

```
let s1 = String::from("hello");  
let s2 = s1;
```



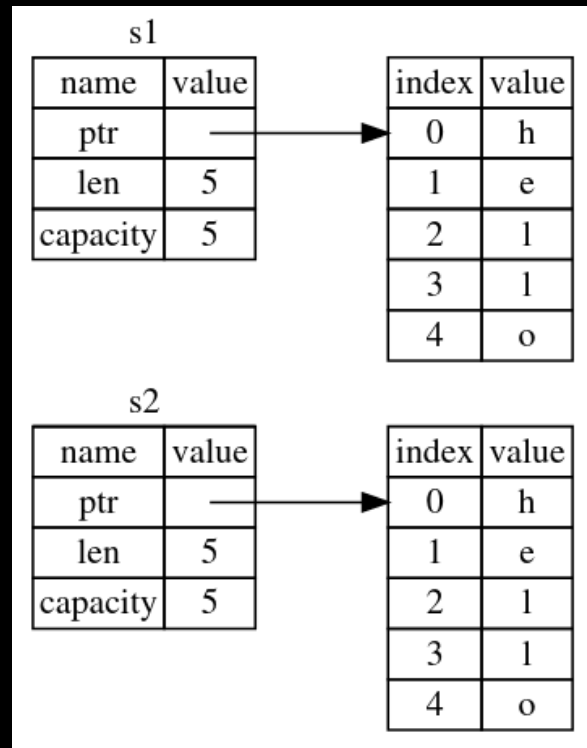

Move and heap-allocated data

```
let s1 = String::from("hello");  
let s2 = s1;
```



Deep copies

```
let s1 = String::from("hello");  
let s2 = s1.clone();
```



Ownership and functions

```
fn main() {  
    let s = String::from("hello");  
    // s comes into scope  
  
    // s's value moves into the function...  
    takes_ownership(s);  
    // ... and so is no longer valid here  
  
    let x = 5; // x comes into scope  
  
    // x would move into the function,  
    // but i32 is Copy, so it's okay to still  
    // use x afterward  
    makes_copy(x);  
} // Here, x goes out of scope, then s.  
// But because s's value was moved, nothing  
// special happens.
```

```
fn takes_ownership(some_string: String) {  
    // some_string comes into scope  
    println!("{}", some_string);  
} // Here, some_string goes out of scope and  
// `drop` is called. The backing  
// memory is freed.  
  
fn makes_copy(some_integer: i32) {  
    // some_integer comes into scope  
    println!("{}", some_integer);  
} // Here, some_integer goes out of scope.  
// Nothing special happens.
```

Ownership and return values

```
fn main() {  
    // gives_ownership moves its return  
    // value into s1  
    let s1 = gives_ownership();  
  
    // s2 comes into scope  
    let s2 = String::from("hello");  
  
    // s2 is moved into takes_and_gives_back,  
    // which also moves its return value into s3  
    let s3 = takes_and_gives_back(s2);  
}  
// Here, s3 goes out of scope and is dropped.  
// s2 was moved, so nothing happens.  
// s1 goes out of scope and is dropped.
```


```
// gives_ownership will move its  
// return value into the function  
// that calls it  
fn gives_ownership() -> String {  
    let some_string = String::from("yours");  
    // some_string comes into scope  
  
    // some_string is returned and  
    // moves out to the calling function  
    some_string  
}  
  
// This function takes a String and returns one  
fn takes_and_gives_back(a_string: String) -> String {  
    // a_string comes into scope  
  
    a_string // a_string is returned and  
    | | | // moves out to the calling function  
}
```

Passing ownership can be annoying

```
fn main() {  
    let s1 = String::from("hello");  
  
    let (s2, len) = calculate_length(s1);  
  
    println!("The length of '{}' is {}.", s2, len);  
}  
  
fn calculate_length(s: String) -> (String, usize) {  
    let length = s.len(); // len() returns the length of a String  
  
    (s, length)  
}
```

Borrows to the rescue!

```
fn main() {  
    let s1 = String::from("hello");  
  
    let len = calculate_length(&s1);  
  
    println!("The length of '{}' is {}.", s1, len);  
}  
  
fn calculate_length(s: &String) -> usize {  
    s.len()  
}
```



Mutable borrows

```
fn main() {  
    let mut s = String::from("hello");  
  
    change(&mut s);  
}  
  
fn change(some_string: &mut String) {  
    some_string.push_str(", world");  
}
```

```
let mut s = String::from("hello");
```

```
let r1 = &mut s;
```

```
let r2 = &mut s;
```

```
// Results in compilation error:  
// println!("{}", {}, r1, r2);
```

```
let mut s = String::from("hello");
```

```
{  
    let r1 = &mut s;  
}
```

```
// r1 goes out of scope here,  
// so we can make a new reference
```

```
let r2 = &mut s;
```

Mutable and shared borrows

```
let mut s = String::from("hello");

let r1 = &s; // no problem
let r2 = &s; // no problem
let r3 = &mut s; // BIG PROBLEM

println!("{}", {}, and {}", r1, r2, r3);
```

```
let mut s = String::from("hello");

let r1 = &s; // no problem
let r2 = &s; // no problem
println!("{}", and {}", r1, r2);
// variables r1 and r2 will not be
// used after this point

let r3 = &mut s; // no problem
println!("{}", r3);
```


Dangling References

```
fn dangle() -> &String {  
    let s = String::from("hello");  
    &s  
}  
  
fn main() {  
    let reference_to_nothing = dangle();  
}
```

Borrow rules

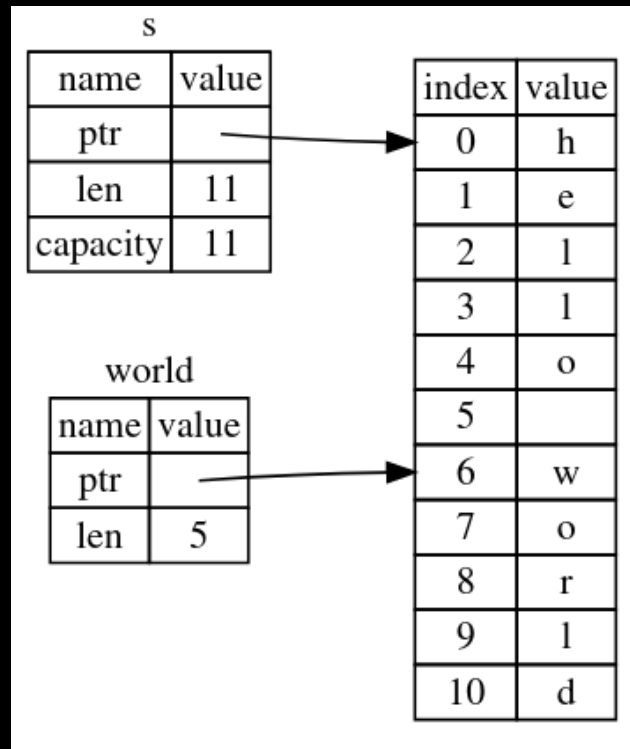
- At any given time, you can have either one mutable reference or any number of immutable references.
- References must always be valid.

Slices

- Let's say we want to write function which returns the first word from string
- It accepts `&String`, but what should it return?
- This can be done using “string slices” `&str`
- Slices can be created by indexing with range syntax `&s[start_idx..end_idx]`

String slices

```
let mut s = String::from("hello");  
  
let s = String::from("hello world");  
  
let hello = &s[0..5];  
let world = &s[6..11];
```



Slices from arrays and vectors

```
fn sum(values: &[u32]) -> u32 {  
    let mut sum = 0;  
    for value in values {  
        sum += value;  
    }  
    sum  
}
```

```
let mut v: Vec<u32> = vec![0, 1, 2, 3];  
v.push(4);  
let a = [0, 1, 2, 3];  
  
let v_sum = sum(&v[1..3]);  
let a_sum = sum(&a[1..2]);
```

Lifetimes

- The main aim of lifetimes is to prevent dangling references
- Borrow checker tracks lifetimes and checks that all borrows are valid
- Borrow checker is based on local reasoning about current function

Implicit lifetimes in code

```
fn main() {  
    let r;                                // -----+-- 'a  
    {                                    // |  
        let x = 5;                       // -+-- 'b |  
        r = &x;                          // | |  
    }..                                  // -+ |  
    println!("r: {}", r);                // |  
}
```

Lifetimes in functions

```
fn longest(x: &str, y: &str) -> &str {  
    if x.len() > y.len() {  
        x  
    } else {  
        y  
    }  
}
```

Lifetime Annotation Syntax

- `&i32` a reference
- `&'a i32` a reference with an explicit lifetime
- `&'a mut i32` a mutable reference with an explicit lifetime
- `&'static i32` a reference with static lifetime

Lifetimes in functions

```
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {  
    if x.len() > y.len() {  
        x  
    } else {  
        y  
    }  
}
```


Lifetimes in types

```
struct LongestStr<'a> {  
    s: &'a str,  
}  
  
fn longest<'a>(x: &'a str, y: &'a str) -> LongestStr<'a> {  
    if x.len() > y.len() {  
        LongestStr { s: x }  
    } else {  
        LongestStr { s: y }  
    }  
}
```

Lifetime ellision

- In simple cases compiler can infer lifetimes
- For example, `fn first_word(x: &str) -> &str`
- Ellison rules:
 - The compiler assigns a lifetime parameter to each parameter that's a reference.
 - If there is exactly one input lifetime parameter, that lifetime is assigned to all output lifetime parameters.
 - If there are multiple input lifetime parameters, but one of them is `&self` or `&mut self` (i.e. it is a method), the lifetime of `self` is assigned to all output lifetime parameters.

Lifetimes in method definitions

```
impl<'a> LongestStr<'a> {  
    fn announce_and_return(&self, announcement: &str) -> &str {  
        println!("Attention please: {}", announcement);  
        self.part  
    }  
}
```

Questions?