

5-1: Multithreading (Theory)

Artem Pavlov, TII, Abu Dhabi, 18.04.2024

Threads

- Common abstraction for independent execution of code
- Execution of threads code can be interleaved or performed in parallel
- Threads in one process work with the common address space
- Each thread usually has its own stack memory

Spawning threads

- Threads can be spawned using `std::thread::spawn` function
- Threads can be configured using `std::thread::Builder` API
- Spawned threads return `JoinHandle` which can be used to wait for spawned thread's termination and retrieval of return result
- Threads can not be killed with `std` APIs!

Spawning scoped threads

- If spawned threads need to borrow non-**static** data, you can use **std::thread::scope** to create “scoped” threads
- Scoped threads involve two lifetimes: **'scope** and **'env** (**'env: 'scope** bound is part of the **Scope** type)
- The **'scope** lifetime represents the lifetime of the scope itself
- The **'env** lifetime represents the lifetime of whatever is borrowed by the scoped threads

Send

- Auto marker trait used for types that can be transferred across thread boundaries
- Commonly used by APIs which deal with sending data across threads
- If **Send** was not implemented for a type, it can be implemented manually using **unsafe**

Sync

- Auto marker trait for types for which it is safe to share references between threads
- In other words, a type `T` is `Sync` if and only if `&T` is `Send`
- `&mut T` is `Sync`, because mutation through `&&mut T` is not possible
- Counter-example: types with interior mutability `Cell` and `RefCell` are not `Sync`
- Atomic types (e.g. `AtomicU32` and `Arc`) are `Sync`

Relation between `Send` and `Sync`

- `&T` is `Send` if and only if `T` is `Sync`
- `&mut T` is `Send` if and only if `T` is `Send`
- `&T` and `&mut T` are `Sync` if and only if `T` is `Sync`
- More in the Rustonomicon:
<https://doc.rust-lang.org/nomicon/send-and-sync.html>

Thread Local Storage

- Can be created using the `thread_local!` Macro
- Creates static with `LocalKey<T>` type
- Only shared (`&T`) references pointing to TLS may be obtained from `LocalKey`
- This restriction can be worked around by using interior mutability, e.g. with `Cell` and `RefCell`

Mutex

- A mutual exclusion primitive useful for protecting shared data
- The protected data can only be accessed through the RAII guards returned from `lock` and `try_lock`
- If during mutex locking code panics, it poisons the mutex
- Common pattern is to use `Arc<Mutex<T>>`

RwLock

- A variation of **Mutex** which supports multiple read locks or exclusive write lock
- The priority policy of the lock is dependent on the underlying operating system's implementation
- Gets poisoned if code panics with acquired write lock

MPSC channels

- Multi-producer, single-consumer FIFO channels
- Defined in the `std::sync::mpsc` module
- `std` provides two variations: unbounded (`mpsc::channel`) and bounded (`mpsc::sync_channel`)
- Channels have two parts: sender and receiver
- Sender types can be cloned

Atomics

- Atomic types provide primitive shared-memory communication between threads
- Std provides the following atomic types:
`AtomicBool`, `AtomicPtr`, `AtomicIsize`, `AtomicUsize`,
`AtomicI8`, `AtomicU16`, `AtomicU32`, etc.
- Atomics can be stored in `statics`

Ordering of atomic operations

- Rust atomics currently follow the C++20 model
- Most atomic operations take Ordering enum
- Ordering has the following variants:
Relaxed, Release, Acquire, AcqRel, SeqCst
- Read more about atomics in the “Rust Atomics and Locks” book: <https://marabos.nl/atomics/>

Case study: `rayon`

- `rayon` is a data-parallelism library for Rust.
- Makes it easy to convert a sequential computation into a parallel one
- Ergonomic: in adaptor-based code you often need only replace `iter` with `par_iter`
- Uses worker stealing thread pool to efficiently implement parallel processing of data
- Example “Computing prime numbers in parallel”:
<https://wannesmalfait.github.io/blog/2024/parallel-primes/>

Questions?